

Parents' Linguistic Alignment Predicts Children's Language Development

Anonymous CogSci submission

Abstract

Children quickly gain enormous linguistic knowledge during early development, in part due to low-level features of their parents' speech. Some posit that parents contribute to their child's language development by calibrating their language according to their child's developmental abilities and needs. Here, we investigate this hypothesis by examining 'alignment' at the level of syntax in a large-scale corpus of parent-child conversations and measuring its association with language development outcomes. To do so, we employ a statistical model of alignment to estimate its presence in our dataset and its impact on language development outcome measures. Our results corroborate previous findings, showing strong alignment for both parents and children; in addition, we demonstrate that parental alignment is a significant predictor of language maturity independent of demographic features, suggesting that parents' tuning has strong ties to their children's language development.

Keywords: Language acquisition; computational modeling; cognitive development

Introduction

Children make vast linguistic strides within their first few years of life, rapidly growing their vocabulary (Mayor & Plunkett, 2010). Some researchers have put forward the linguistic tuning hypothesis, arguing that parents bolster their child's early language learning by calibrating the complexity of their speech to the particular abilities and needs of their children (Montag & MacDonald, 2015; Snow, 1972; Thiessen, Hill, & Saffran, 2005). The idea is intuitive, but it is unclear at what level tuning occurs (Hayes & Ahrens, 1988; Sokolov, 1993; Spivey & Dale, 2006) and how overt it is (Brown & Hanlon, 1970; Chouinard & Clark, 2003; Hirsh-Pasek, Treiman, & Schneiderman, 1984).

A parallel yet complementary vein of language development research investigates the presence of low-level cues in parental speech and their influence on child language learning. From this research, we know that child-directed speech contains features that facilitate language learning, and that more exposure tends to result in better outcomes (Cameron-Faulkner, Lieven, & Tomasello, 2003; Weisleder & Fernald, 2013). Related, caregivers from families of high socioeconomic status (SES) tend to converse more with their children than their lower SES counterparts, and these increases are associated with improved development outcomes such as vocabulary size and school performance (Hoff, 2003; Walker, Greenwood, Hart, & Carta, 1994). Moreover, differences

in SES-based language development are largely explained by low-level features of parental child-directed speech such as lexical diversity and sentence complexity (Hoff-Ginsberg, 1998; Rowe, 2008). So, given that granular aspects of parent speech can have substantial effects on a child's language development, it may be that linguistic tuning occurs at this level in subtle ways, particularly when it comes to non-content words (i.e., words that are not central to the topic of discussion.)

This idea of assessing the direct impact of a parent's usage of non-content word on language development relates to linguistic alignment, a phenomenon whereby conversational partners tend to align aspects of their communicative style and content according to various external influences e.g., (Pennebaker, Booth, Boyd, & Francis, 2015). Yurovsky, Doyle, & Frank (2016) investigates linguistic alignment in CHILDES (MacWhinney, 2000), a natural language corpus of conversation between parents and children to assess whether tuning occurs at the level of syntactic categories. They find that syntactic alignment does occur between both parents and children; moreover, parents align less over time, suggesting that the relationship their speech shares with their child's changes as a function of development. These results present a powerful proof of concept that syntactic alignment exists between parents and children and changes over time, but it remains unclear whether alignment bears any sort of concrete, impactful relationship to language development.

Here, we extend Yurovsky, Doyle, and Frank's (2016) model by applying it to the Language Development Project (LDP) (Goldin-Meadow et al., 2014), a corpus of ecological conversations between parents and their children over time, collected from a socioeconomically diverse sample of parent-child dyads. Moreover, we use alignment estimates alongside demographic information to predict language outcome measures, strengthening the linguistic tuning hypothesis by concretely showing how parents' sensitivity to their child's linguistic needs and abilities impacts their development.

Model

The linguistic tuning hypothesis predicts that parents will calibrate their language in part by assessing their child's needs and abilities. So, we predict that parents will exhibit high alignment to their young children, but will reduce their alignment as their children mature. To test this prediction, we

employ an extended version of the Hierarchical Alignment Model (HAM) implemented in Yurovsky et al. (2016) which both estimates the impact of a speaker’s use of syntactic categories on their conversational partner’s usage and uses these alignment estimates to predict language outcome scores.

At base, the model predicts, for each utterance, whether the speaker will produce a word from a given syntactic category. This prediction is generated by two factors: the speaker’s baseline propensity towards using that category and the speaker’s tendency to align, producing words from a category just used by their partner. In the model, the primary computation mimics a standard logistic regression - the production of a syntactic category within an utterance is treated as a binary outcome variable impacted by a linear combination of predictor variables (here, baseline usage and alignment.) The model’s hierarchical structure then allows the estimates of baseline usage and alignment effects to be pooled across individual speakers and syntactic categories in a way that ensures statistical robustness.

The model used here then incorporates these alignment and baseline usage estimates as predictors in a linear regression model of PPVT (Dunn & Dunn, 1997). At this stage, PPVT is estimated as a linear combination of predictors reflecting alignment and baseline usage estimate for both parents and children alongside other predictors representing demographic variables (e.g., child’s gender, mother’s education) and the child’s age.

| | category | examples |
|----|----------|-------------------|
| 1 | article | a, alot |
| 2 | certain | altogether, must |
| 3 | conj | but, or |
| 4 | discrep | wanted, hoped |
| 5 | excl | whether, not |
| 6 | i | i’m, i |
| 7 | incl | both, around |
| 8 | ipron | thatd, whats |
| 9 | negate | needn’t, oughtn’t |
| 10 | preps | at, to |
| 11 | quant | series, every |
| 12 | tentat | anyhow, most |
| 13 | we | we’d, lets |
| 14 | you | youd, y’all |

Table 1: LIWC Categories with example words.

Model Details

The structure of the model used here greatly resembles that used in Yurovsky et al. (2016), in that it operates over utterances represented as binary vectors, with indices indicating the presence or absence of each of the 14 LIWC categories (Table 1). The probability of producing each category in each utterance is predicted from two parameters: the speaker’s baseline usage of that syntactic category (η_c^{base}), and the change in that speaker’s baseline as a function of in-

teracting with the listener (η_c^{align}). So, for a given syntactic category c , for replies to utterances that don’t contain c , the production parameter for that category is computed by applying the inverse logit function to the appropriate baseline log odds:

$$P(Production_c) = \text{logit}^{-1}(\eta_c^{base})$$

Alternatively, replies to utterances that *do* contain c , the parameter computation takes into account the sum of the baseline and alignment log odds:

$$P(Production_c) = \text{logit}^{-1}(\eta_c^{base} + \eta_c^{align})$$

To accommodate the variance in production across the LIWC categories, each baseline usage parameter was drawn from an uninformative prior ($\eta_c^{base} \sim \text{Uniform}(-5, 5)$); alignment parameters were regularized towards 0 by way of implementing a conservative prior ($\eta_c^{align} \sim \text{Normal}(0, .25)$).

All parameters were estimated hierarchically, which allows intelligent pooling of data across participants in the dataset. To start, each subpopulation (i.e., parents vs. children) obtained an estimate. Then, every speaker had an alignment estimate drawn from their appropriate subpopulation (e.g., if Speaker 22 is a child, their alignment estimate is drawn from the estimate for children overall.) Category-level alignment estimates were then drawn for each speaker (e.g., the alignment estimate for Speaker 22’s usage of determiners is drawn from Speaker 22’s alignment estimate.) The order was flipped for baseline estimates in order to better reflect empirical baseline usages across syntactic categories; subpopulation estimates produced category-level estimates, which then produced speaker-level estimates. As in Yurovsky et al. (2016), we also include parameters that allow baseline and alignment probabilities to change linearly over time (β and α respectively).

Next, we extend the model used in Yurovsky et al. (2016) by using estimated alignment (i.e., η parameters) to predict PPVT scores, a widely used inventory for tracking language development (Dunn & Dunn, 1997). To do so, we implement a regression model where PPVT scores are modeled as linear combinations of various predictor variables. These predictor variables included the child’s age, alignment parameter estimates for the child and their parent, the mother’s education, the child’s gender, as well as interaction effects for all variables with age. Error variance for the model (σ) was also estimated.

The model implemented here then serves two purposes: (1) It extends the analysis of Yurovsky et al. (2016) to a new dataset, aiming to replicate previous findings in a more diverse and representative sample, and (2) It incorporates alignment estimates in a predictive model of early language outcomes, serving to test the hypothesis that alignment has significant effects on language development over and above demographic features. To be specific, we hope to replicate non-zero estimates for η parameters (demonstrating that alignment between parents and children exists across datasets), positive β for children (showing that children increase their

baseline usage of categories over time), and negative α for parents (showing that parents decrease their alignment as their children age.) If the PPVT model estimates for parameters corresponding to the main or interaction effects of alignment are non-zero in the presence of demographic variables, we can infer that alignment has an relationship with language development independent of features like socioeconomic status, bolstering the linguistic tuning hypothesis.

Analysis

Data

Conversations between parents and their children were drawn from the Language Development Project Corpus (Goldin-Meadow et al., 2014). Participants in the project were video-recorded in their homes for ~ 90 minutes every four months starting when the child was 14-months and ending at 58-months. Participants were selected in order to produce a diverse sample demographically representative of the broader Chicagoland area.

We selected for analysis all children who were typically developing and completed at least 10 of the 12 planned recording sessions. Our sample consisted of 59 target children, 28 of whom were girls, 12 were Black and 6 were Multiracial. Children were also socio-economically diverse, as measured by mother’s education: 2 mothers had some highschool education, 7 had a highschool degree, 10 had some college or trade school, 19 had a highschool degree or GED, 19 had college degrees, and 21 had advanced degrees.

Following Yurovsky et al. (2016), successive utterances from a speaker within a transcript were concatenated into a single utterance. Individual utterances were then transformed into binary vectors with indices indicating the presence or absence of each of the 14 LIWC categories. This pre-processing turned every transcript into a speaker-reply format: each utterance within a transcript was both a reply to the preceding utterance and a message to the next one.

Each transcript was then compressed, yielding 4 numbers for each LIWC category. For a pair of speakers A and B in a transcript, for each LIWC category, we computed the number of utterances from A to B containing the category (N^{align}), the number of utterances from A to B not containing the category (N^{base}), the number of utterances containing the category responding to an utterance containing the category (C^{align}), and the number of utterances containing the category responding to an utterance not containing the category (C^{base}). Aggregating in this way provided the platform for the model’s sampling - for each transcript, C^{base} and C^{align} were drawn from Binomial distributions parameterized by N^{base} and N^{align} chances respectively, with probabilities computed via the logistic regression models outlined above.

Sampling was performed using Stan, a probabilistic programming language that implements Hamiltonian Monte Carlo sampling methods (Carpenter et al., 2017). Posterior distributions for each parameter in the model were estimated using




Figure 1: Posterior parameter estimates for alignment (η), developmental change in alignment (α), and developmental change in baseline function word production (β) for both parents and children, as well estimated alignment between parents for a baseline. Bars indicate means, error-bars indicate 95% highest posterior density intervals

500 iterations.¹

Results

Alignment estimates (η^{align}) for parents and children were both estimated above zero, corroborating the findings of Yurovsky et al. (2016) in showing that both groups exhibit alignment (Figure 1). We also replicate the finding that parents appear to align more to their children than children align to their parents.

The model estimates developmental baseline changes (β) at approximately zero for parents, but significantly above zero for children, replicating previous findings. Alignment is estimated as having no significant age effect (α) in this dataset, failing to replicate an earlier finding that parents tend to exhibit decreased alignment over their child’s development. However, the mean estimate is trending negatively, suggesting this may be a function of the data’s limited scope (Figure 2).

The mean estimates for PPVT predictors are presented in Figure 3. PPVT is positively associated with the age of the child and their being female. Moreover, female children tend to have a decreased age effect on PPVT; this means that for female children, the expected increase in PPVT over time is smaller. Mother’s education is negatively associated with PPVT, but has a slight positive age effect. Alongside these demographic effects, we see robust alignment effects on PPVT: child and parental alignment are both associated with increased PPVT, but with decreased age effects.

Discussion

In an effort to understand and investigate how children rapidly acquire language, some argue that the language parents produce to their children is somehow calibrated to the child’s

¹Data and code available on Github; will be made available at time of acceptance.




Figure 2: Model-estimated changes in linguistic alignment over development. Points indicate the mean of the posterior distribution, shaded regions indicate 68% highest probability density intervals, equivalent to one standard deviation

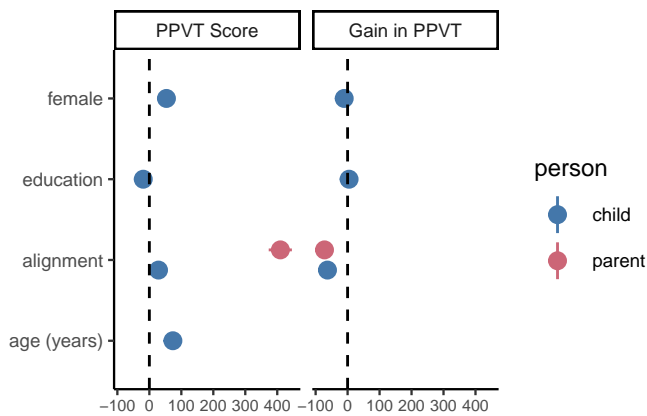


Figure 3: Posterior parameter estimates for demographic effects on PPVT scores. Points indicate means, error-bars indicate 95% highest posterior density intervals

particular needs and abilities (Snow, 1972). While the idea is theoretically compelling, empirical work has produced mixed results, with strong results in favor of (Chouinard & Clark, 2003; Hirsh-Pasek et al., 1984) and against (Brown & Hanlon, 1970; Hayes & Ahrens, 1988).

However, much of this prior work investigates tuning as an overt effort on behalf of parents or tuning with respect to content words, with less examining the potential role of low-level syntactic influence (Hoff, 2003). Yurovsky et al. (2016) presents just such an examination, demonstrating using Bayesian hierarchical modeling that parents align to their children according to their particular language usage at the level of syntactic categories. This paper extends their model by applying it to a new socioeconomically diverse sample of families (Goldin-Meadow et al., 2014) and leveraging the model's alignment estimates to predict language development outcomes.

The analysis presented here replicates the findings of Yurovsky et al. (2016), showing strong alignment effects for both parents and their children, a substantial age effect for baseline useage in children, and a trending negative effect of age on alignment for parents. Moreover, we demonstrate that these alignment estimates have substantial power in predicting language development outcomes, even in the presence of demographic features such as gender and socioeconomic status. We corroborate previous findings that female children tend to have higher PPVT scores that diminish over time (Kaushanskaya, Gross, & Buac, 2013; Lange, Euler, & Zaretsky, 2016). We conflict with other findings that positively associate child PPVT scores with mother's education (Di Cesare, Sabates, & Lewin, 2013; Schady, 2011); this may be due to idiosyncracies of our dataset, including its limited size.

We show that parental alignment is associated with a relatively large boost in overall PPVT scores, but with a negative age effect. The negative age effect may source from a ceiling on PPVT - children with higher scores may simply have less ground to cover. Nevertheless, these results are consistent with a concrete effect of parental alignment on language development, and the linguistic tuning hypothesis more broadly. A similar story is evident from child alignment estimates: alignment has a small association with overall PPVT score and an age effect comparable to parental alignment. Here there may be a confound with childrens' baseline language production, in that children with lower production will have lower PPVT and diminished alignment as a result; future work should assess this interaction to better isolate the effects of alignment.

Overall, these results show that parental alignment is a robust effect that appears to have a relationship with childrens' language development independent of demographic correlates, serving to further the linguistic tuning hypothesis.

References

- Brown, R. W., & Hanlon, C. (1970). Derivational Complexity and Order of Acquisition in Child Speech. In J. Hayes (Ed.), *Cognition and the development of language* (pp. 11–53). New York.
- Cameron-Faulkner, T., Lieven, E., & Tomasello, M. (2003). A construction based analysis of child directed speech. *Cognitive Science*, 27(6), 843–873.
- Carpenter, B., Gelman, A., Hoffman, M. D., Lee, D., Goodrich, B., Betancourt, M., ... Riddell, A. (2017). Stan : A Probabilistic Programming Language. *Journal of Statistical Software*, 76(1).
- Chouinard, M. M., & Clark, E. V. (2003). Adult reformulations of child errors as negative evidence. *Journal of Child Language*, 30(3), 637–669.
- Di Cesare, M., Sabates, R., & Lewin, K. M. (2013). A double prevention: how maternal education can affect maternal mental health, child health and child cognitive development. *Longitudinal and Life Course Studies*, 4.
- Dunn, L. M., & Dunn, L. M. (1997). Peabody Picture Vocabulary Test—Third Edition. *PsycTESTS Dataset*.
- Goldin-Meadow, S., Levine, S. C., Hedges, L. V., Huttenlocher, J., Raudenbush, S. W., & Small, S. L. (2014). New evidence about language and cognitive development based on a longitudinal study: Hypotheses for intervention. *American Psychologist*, 69(6), 588–599.
- Hayes, D. P., & Ahrens, M. G. (1988). Vocabulary simplification for children: a special case of 'motherese'? *Journal of Child Language*, 15(2), 395–410.
- Hirsh-Pasek, K., Treiman, R., & Schneiderman, M. (1984). Brown & Hanlon revisited: mothers' sensitivity to ungrammatical forms. *Journal of Child Language*, 11(01), 81–88.
- Hoff, E. (2003). The specificity of environmental influence: Socioeconomic status affects early vocabulary development via maternal speech. *Child Development*, 74(5), 1368–1378.
- Hoff-Ginsberg, E. (1998). The relation of birth order and socioeconomic status to children's language experience and language development. *Applied Psycholinguistics*, 19, 603–629.
- Kaushanskaya, M., Gross, M., & Buac, M. (2013). Gender differences in child word learning. *Learning and Individual Differences*, 27, 82–89.
- Lange, B. P., Euler, H. A., & Zaretsky, E. (2016). Sex differences in language competence of 3- to 6-year-old children. *Applied Psycholinguistics*, 37(06), 1417–1438.
- MacWhinney, B. (2000). The CHILDES Project. *Computational Linguistics*, 26(4), 657–657.
- Mayor, J., & Plunkett, K. (2010). A statistical estimate of infant and toddler vocabulary size from CDI analysis. *Developmental Science*, 14(4), 769–785.
- Montag, J. L., & MacDonald, M. C. (2015). Text exposure predicts spoken production of complex sentences in 8- and 12-year-old children and adults. *Journal of Experimental Psychology: General*, 144(2), 447–468.
- Pennebaker, J. W., Booth, R. J., Boyd, R. L., & Francis, M. E. (2015). Linguistic inquiry and word count: LIWC2015. Austin, TX: Pennebaker Conglomerates.
- Rowe, M. L. (2008). Child-directed speech: relation to socioeconomic status, knowledge of child development and child vocabulary skill. *Journal of Child Language*, 35(01), 185–205.
- Schady, N. (2011). Parents' education, mothers' vocabulary, and cognitive development in early childhood: longitudinal evidence from Ecuador. *American Journal of Public Health*, 101(12), 2299–2307.
- Snow, C. E. (1972). Mothers' Speech to Children Learning Language. *Child Development*, 43(2), 549–565.
- Sokolov, J. L. (1993). A local contingency analysis of the fine-tuning hypothesis. *Developmental Psychology*, 29(6), 1008–1023.
- Spivey, M. J., & Dale, R. (2006). Continuous Dynamics in Real-Time Cognition. *Current Directions in Psychological Science*, 15(5), 207–211.
- Thiessen, E. D., Hill, E. A., & Saffran, J. R. (2005). Infant-Directed Speech Facilitates Word Segmentation. *Infancy*, 7(1), 53–71.
- Walker, D., Greenwood, C., Hart, B., & Carta, J. (1994). Prediction of School Outcomes Based on Early Language Production and Socioeconomic Factors. *Children and Poverty*, 65(2), 606–621.
- Weisleder, A., & Fernald, A. (2013). Talking to Children Matters. *Psychological Science*, 24(11), 2143–2152.
- Yurovsky, D., Doyle, G., & Frank, M. C. (2016). Linguistic input is tuned to children's developmental level. In *Proceedings of the annual meeting of the cognitive science society* (pp. 2093–2098).